

# CSM CS61C Note #3: RISC-V Calling Conventions and Instruction Formats

Anthony Han

## 1 RISC-V Calling Conventions

In RISC-V, registers are divided into **caller-saved registers** and **callee-saved registers**. Programmers are expected to make sure that before and after a function call, the value stored in callee-saved registers stay the same. This is the RISC-V **calling convention**.

| Register | Name    | Description                        | Saver  |
|----------|---------|------------------------------------|--------|
| x0       | zero    | Always 0                           | N/A    |
| x1       | ra      | Return address                     | Caller |
| x2       | sp      | Stack pointer                      | Callee |
| x3       | gp      | Global pointer                     | N/A    |
| x4       | tp      | Thread pointer                     | N/A    |
| x5 – 7   | t0 – 2  | Temporary                          | Caller |
| x8       | s0 / fp | Saved register / Frame pointer     | Callee |
| x9       | s1      | Saved register                     | Callee |
| x10 – 17 | a0 – 7  | Function arguments / Return values | Caller |
| x18 – 27 | s2 – 11 | Saved registers                    | Callee |
| x28 – 31 | t3 – 6  | Temporaries                        | Caller |

You can also find this table on the RISC-V green sheet. Note that:

- Return address ra stores where to return to in a nested procedure.
- Stack pointer sp points to the bottom of call frame. Convention is grow stack down from high to low addresses: push decrements sp, pop increments sp.
- If there are more than 8 arguments, put on stack:

```
1 | sw t0, 0(sp) # Argument #9
2 | sw t1, 4(sp) # Argument #10
3 | # You need to memorize where they are stored.
```

As a result of this calling convention, within each function, we tend to have a prologue and epilogue that looks like

```
1 | func:
2 |     # Prologue
3 |     addi sp, sp, -8 # Adjust stack for 2 items
4 |     sw ra, 4(sp)   # Save ra for use afterwards
5 |     sw s0, 0(sp)   # Save s0 for use afterwards
6 |
7 |     # Do stuff
```

```

8 ||
9   # Epilogue
10  lw s0, 0(sp)    # Restore register s0 for caller
11  lw ra, 4(sp)    # Restore register ra for return
12  addi sp, sp, 8   # Adjust stack to delete 2 items
13  jr ra           # Jump back to calling routine


```

## 2 RISC-V Instruction Formats

To let the computer execute the assembly we write, we need to turn the instructions into raw bits. Last time around, we classified RISC-V instructions into five categories by what these instructions do. This time, we are going to classify these instructions by how they are encoded in machine code. There are some similarities, yet there are differences.

### RISC-V Instruction Format - R Format

This corresponds to the register instructions we discussed last time.



- The **opcode** field of an instruction partially specifies which instruction it is. It is **equal to (0110011)<sub>2</sub> for all R-type instructions**.
- The **funct7** and **funct3** fields, when combined with the opcode, describe the operation to be performed.
- Each register field (**rs1**, **rs2**, **rd**) holds a 5-bit unsigned integer (0–31) corresponding to a register (x0–31).
  - **rs1** (Source register #1): specifies register containing first operand.
  - **rs2** (Source register #2): specifies register containing second operand.
  - **rd** (Destination register): specifies register which will receive result of computation.

- R format example:

```
1 || add x18, x19, x10
```

|         |        |        |     |       |            |
|---------|--------|--------|-----|-------|------------|
| 0000000 | 01010  | 10011  | 000 | 10010 | 0110011    |
| ADD     | rs2=10 | rs1=19 | ADD | rd=18 | Reg-Reg OP |


- All R-type instructions:

|         |     |     |     |    |         |      |
|---------|-----|-----|-----|----|---------|------|
| 0000000 | rs2 | rs1 | 000 | rd | 0110011 | ADD  |
| 0100000 | rs2 | rs1 | 000 | rd | 0110011 | SUB  |
| 0000000 | rs2 | rs1 | 001 | rd | 0110011 | SLL  |
| 0000000 | rs2 | rs1 | 010 | rd | 0110011 | SLT  |
| 0000000 | rs2 | rs1 | 011 | rd | 0110011 | SLTU |
| 0000000 | rs2 | rs1 | 100 | rd | 0110011 | XOR  |
| 0000000 | rs2 | rs1 | 101 | rd | 0110011 | SRL  |
| 0100000 | rs2 | rs1 | 101 | rd | 0110011 | SRA  |
| 0000000 | rs2 | rs1 | 110 | rd | 0110011 | OR   |
| 0000000 | rs2 | rs1 | 111 | rd | 0110011 | AND  |

## RISC-V Instruction Format - I Format

Ideally, RISC-V only uses one instruction format; this is not possible, however: 5b field only represents numbers up to the value 31.

Therefore, to represent the immediate instructions, we define a format most consistent with R-format which uses at most 2 registers (one source, one destination):



- Here, the rs2 and funct7 fields are replaced by a 12-bit signed immediate, **imm[11:0]**. The remaining fields stay the same.
- Immediate is **always sign-extended to 32b** before used in an arithmetic operation.
- Some I-type instructions:

|           |       |     |     |    |         |       |
|-----------|-------|-----|-----|----|---------|-------|
| imm[11:0] |       | rs1 | 000 | rd | 0010011 | ADDI  |
| imm[11:0] |       | rs1 | 010 | rd | 0010011 | SLTI  |
| imm[11:0] |       | rs1 | 011 | rd | 0010011 | SLTIU |
| imm[11:0] |       | rs1 | 100 | rd | 0010011 | XORI  |
| imm[11:0] |       | rs1 | 110 | rd | 0010011 | ORI   |
| imm[11:0] |       | rs1 | 111 | rd | 0010011 | ANDI  |
| 0000000   | shamt | rs1 | 001 | rd | 0010011 | SLLI  |
| 0000000   | shamt | rs1 | 101 | rd | 0010011 | SRLI  |
| 0100000   | shamt | rs1 | 101 | rd | 0010011 | SRAI  |


For logical shifts, we note that:

- One of the higher-order immediate bits is used to distinguish SRLI from SRAI.
- We only use the lower 5 bits of the immediate value for shift amount, since we can only shift by 0–31 bit positions.
- **Load instructions are also I-type.** We use the funct3 field to encode the size and signedness of data we want to load.

|           |  |     |     |    |         |     |
|-----------|--|-----|-----|----|---------|-----|
| imm[11:0] |  | rs1 | 000 | rd | 0000011 | LB  |
| imm[11:0] |  | rs1 | 001 | rd | 0000011 | LH  |
| imm[11:0] |  | rs1 | 010 | rd | 0000011 | LW  |
| imm[11:0] |  | rs1 | 100 | rd | 0000011 | LBU |
| imm[11:0] |  | rs1 | 101 | rd | 0000011 | LHU |

## RISC-V Instruction Format - S Format

This corresponds to the store instructions we discussed last time.



- We use rs1 for the base memory address, rs2 for the data to be stored.
- All S-type instructions:

|           |     |     |     |          |         |    |
|-----------|-----|-----|-----|----------|---------|----|
| imm[11:5] | rs2 | rs1 | 000 | imm[4:0] | 0100011 | SB |
| imm[11:5] | rs2 | rs1 | 001 | imm[4:0] | 0100011 | SH |
| imm[11:5] | rs2 | rs1 | 010 | imm[4:0] | 0100011 | SW |

## RISC-V Instruction Format - SB Format

To encode the branch instructions, we need a way to represent the labels. In RISC-V, this is done through **PC-relative addressing**.

|         |           |       |       |        |          |         |        |   |   |
|---------|-----------|-------|-------|--------|----------|---------|--------|---|---|
| 31      | 30        | 25 24 | 20 19 | 15 14  | 12 11    | 8       | 7      | 6 | 0 |
| imm[12] | imm[10:5] | rs2   | rs1   | funct3 | imm[4:1] | imm[11] | opcode |   |   |
| 1       | 6         | 5     | 5     | 3      | 4        | 1       | 7      |   |   |

- We use the **immediate** field as a **two's complement offset relative to the PC**. The remaining fields are identical to those of S-type instructions.
- Since RISC-V instructions are at least 16b, the byte offset between instructions must be even, and **the least significant bit can be omitted**; and we can encode imm[12:1] instead of imm[11:0].
- All SB-type instructions:

|              |     |     |     |             |         |      |
|--------------|-----|-----|-----|-------------|---------|------|
| imm[12:10:5] | rs2 | rs1 | 000 | imm[4:1:11] | 1100011 | BEQ  |
| imm[12:10:5] | rs2 | rs1 | 001 | imm[4:1:11] | 1100011 | BNE  |
| imm[12:10:5] | rs2 | rs1 | 100 | imm[4:1:11] | 1100011 | BLT  |
| imm[12:10:5] | rs2 | rs1 | 101 | imm[4:1:11] | 1100011 | BGE  |
| imm[12:10:5] | rs2 | rs1 | 110 | imm[4:1:11] | 1100011 | BLTU |
| imm[12:10:5] | rs2 | rs1 | 111 | imm[4:1:11] | 1100011 | BGEU |

- If the offset is more than  $2^{12}$  bytes away from the branch instruction, we can use a jump instruction instead:

```

1 || beq x10, x0, far
2 || # next instr
3 ||
4 || # CHANGE TO
5 || bne x10, x0, next
6 || j far
7 || next: # next instr

```

## RISC-V Instruction Format - U Format

As we can see, the immediate in I-type instruction is fairly small. What if we want to express large immediates? Hence the upper immediate instruction. In the case of U-type instructions, the immediate is the **upper 20 bits** of a 32b instruction word.

- **LUI - Load Upper Immediate**
  - Writes the **upper 20 bits** of the destination with the immediate value, and **clears the lower 12 bits**.
  - Together with an **addi** to set low 12 bits, we can create any 32b value in a register using two instructions.

```

1 || lui x10, 0x87654 # x10 = 0x87654000
2 || addi x10, x10, 0x321 # x10 = 0x87654321

```

– A corner case: if we want to set 0xDEADBEEF...

```


1 | lui x10, 0xDEADB # x10 = 0xDEADB000
2 | addi x10, x10, 0xEEF # x10 = 0xDEADAEEF
3 | # Wrong result because immediate in addi is sign-extended
4 | # Instead, pre-increment the value in upper 20 bits:
5 | lui x10, 0xDEADC # x10 = 0xDEADC000
6 | addi x10, x10, 0xEEF # x10 = 0xDEADBEEF
7 |
8 | # In fact, the assembler pseudo-instruction handles it all:
9 | li x10, 0xDEADBEEF # Creates two instructions

```

- **AUIPC - Add Upper Immediate to PC**

## RISC-V Instruction Format - UJ Format

JAL's syntax is fairly different, and so we put it in its own category. Since JAL takes in a label, we are doing a **PC-relative jump**, just like in **SB-type instructions**.



- Saves  $PC + 4$  in register **rd** as return address;
- Sets  $PC = PC + \text{offset}$ .


```

1 | # j - pseudo-instruction, sets rd = x0
2 | j label # == jal x0, label # Discard return address
3 |
4 | # Call function within 2**18 instructions of PC
5 | jal ra, FuncName

```

## JALR Instruction - I Format

Note that in RISC-V, jalr is an I-format instruction:



- Writes  $PC + 4$  to **rd** as return address;
- Sets  $PC = rs + \text{immediate}$ .
- Uses same immediates as arithmetic and loads - **no multiplication by 2 bytes**.

```

1 | # jr - pseudo-instruction, sets rd = x0, imm = 0
2 | jr ra # == jalr x0, ra, 0
3 | # ret - pseudo-instruction
4 | ret # == jr ra
5 |
6 | # Call function at any 32-bit absolute address
7 | lui x1, <hi20bits>
8 | jalr ra, x1, <lo12bits>

```

```
9  |
10 | # Jump PC-relative with 32-bit offset
11 | auipc x1, <hi20bits>
12 |     # Adds upper immediate value to and places result in x1
13 | jalr x0, x1, <lo12bits>
14 |     # Same sign extension trick needed as LUI
```