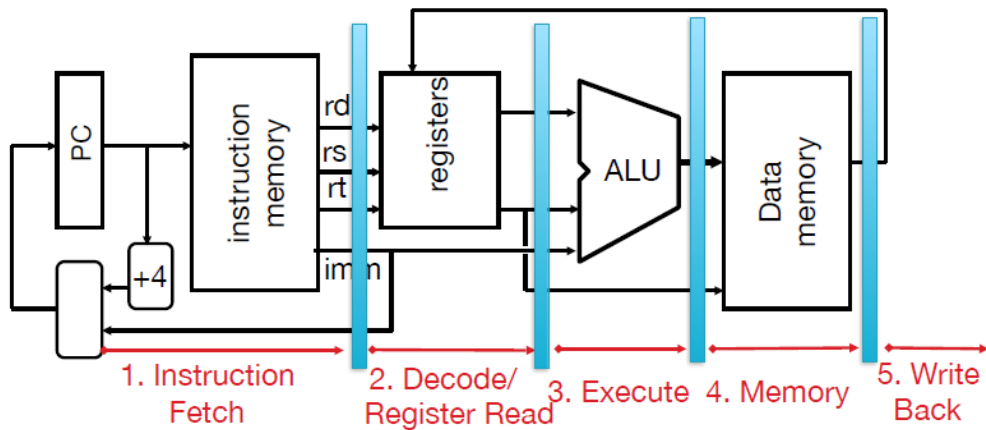


CSM CS61C Note #6: Pipelining, Pipelining Hazards

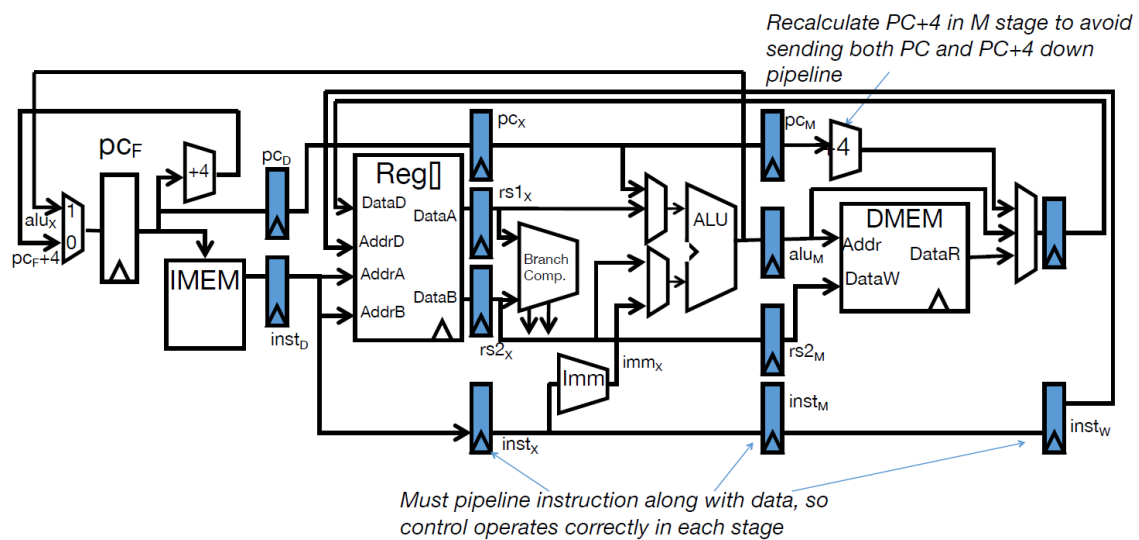
Anthony Han

1 Pipelined RISC-V Datapath







In a single-cycle datapath, only one component is doing the computations. To increase the clock frequency, we pipeline the circuit so that multiple components are in use at the same time, using **registers between stages to hold information produced in previous cycles**:




And we have our pipelined 5-stage RISC-V processor:



- **Increases throughput** - multiple tasks operating simultaneously using different resources.
- **Never improves latency; sometimes worse.**
- **Potential speedup equals the number of pipeline stages.**
- **Time to “fill” and “drain” pipeline reduces speedup;** approaches potential speedup as number of instructions increases.

Phase	Pictogram	t_{step} Serial	t_{cycle} Pipelined
Instruction Fetch		200 ps	200 ps
Reg Read		100 ps	200 ps
ALU		200 ps	200 ps
Memory		200 ps	200 ps
Register Write		100 ps	200 ps
$t_{instruction}$		800 ps	1000 ps



Starting from the fifth cycle, all five stages will be at work simultaneously.

2 Pipelining Hazards

A **hazard** is a situation that **prevents us from starting the next instruction in the next clock cycle.**


Structural Hazard

- Two or more **instructions** in the pipeline **race for access to a single physical resource.**
- Two possible solutions:
 - Solution 1: Instructions take turns to use resource; **stall some instructions.**
 - Solution 2: **Add more hardware.** We can always solve a structural hazard by adding more hardware.
- **Regfile Structural Hazards**
 - Each instruction can read up to two operands in decode stage and write one value in writeback stage.

- We can **avoid structural hazard** by having separate “ports”: two independent read ports and one independent write port.
- **Two reads and one write can happen simultaneously in one cycle.**

- **Memory Structural Hazards**

- Instruction and data memory used simultaneously.
- We could **just use two separate memories.**
- Instruction and data caches:




- RISC ISAs are **designed to avoid structural hazards** - at most 1 memory access/instruction.

Data Hazard


- **Data dependency between instructions.**
- **Data Hazard: Register Access**

- Although we have separate ports for read and write, but when two pipelined instructions **simultaneously read and write to a register**, the read may have ambiguous behavior.



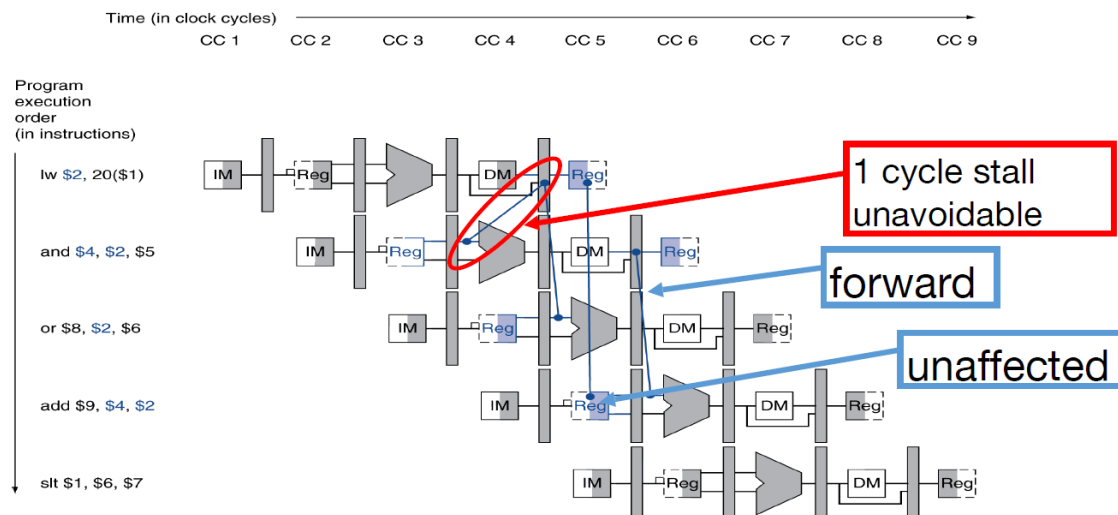
- We could **exploit high speed of register file** (≈ 100 ps) to fit everything in a clock cycle (≈ 200 ps): **WB updates value first, then ID reads new value in.**
- We note it might not always be possible to write then read in the same cycle, especially in high-frequency designs.

- **Data Hazard: ALU Result**



Without some fix, sub and or will give wrong results as these lines are dependent on the s0 value generated in the add instruction!

• Data Hazard: Load Word



We note that in the example above, x2 is used straight after it being loaded from memory.

- Slot after a load is called a **load delay slot**. If that instruction uses the result of the load, then the hardware will have to stall for one cycle.
- To deal with this, we could
 - * Put unrelated instruction into load delay slot to avoid use of load result in the next instruction.
 - * Consider the RISC-V code for C code:

```

1 || A[3] = A[0] + A[1];
2 || A[4] = A[0] + A[2];

```

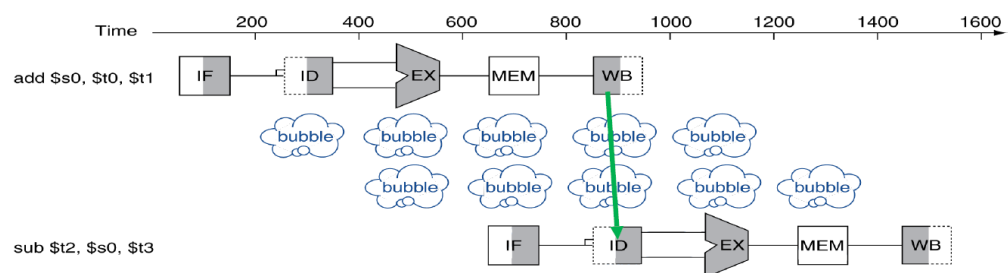
1	# Original Order - 9 Cycles	1	# Alternative - 7 Cycles
2	lw t1, 0(t0)	2	lw t1, 0(t0)
3	lw t2, 4(t0)	3	lw t2, 4(t0)
4	add t3, t1, t2 # Stall!	4	lw t4, 8(t0)
5	sw t3, 12(t0)	5	add t3, t1, t2 # Move!
6	lw t4, 8(t0)	6	sw t3, 12(t0)
7	add t5, t1, t4 # Stall!	7	add t5, t1, t4 # Move!
8	sw t5, 16(t0)	8	sw t5, 16(t0)

- Two possible solutions:

– Solution 1: **Stalling**

- * In the following code block, an instruction depends on the result from a previous instruction:


```
1 | add s0, t0, t1
2 | sub t2, s0, t3
```



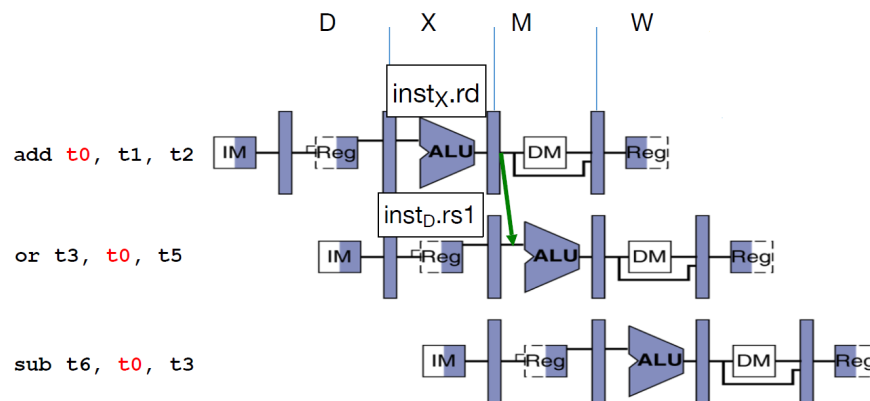
Here, we **fill in NOPs** between the two instructions; the affected pipeline stages would simply do “nothing.”

- * **Stalls reduce performance**; but are **sometimes required to get correct results**.
 - * **Compiler can rearrange code or insert NOPs (writes to register x0) to avoid hazards and stalls.** This, however, requires knowledge of the pipeline structure.
- Solution 2: **Forwarding** (or Bypassing)

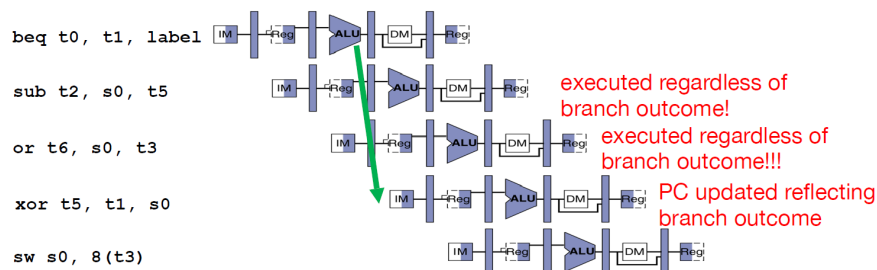
- * **Grab operand from pipeline stage when it is computed.**
 - Do not wait for it to be stored in a register.
 - **Requires extra connections in the datapath.**



- * To detect need for forwarding, we **compare the destination register of older instructions in the pipeline with the source registers of new instruction in decode stage**.



Control Hazards



- If branch is not taken, then instructions fetched sequentially after branch should remain.
- **If branch is taken, then we need to flush/kill incorrect instructions from pipeline by converting them to NOPs.** Thus every taken branch in our simple pipeline costs 2 dead cycles.
- To improve performance, use “**branch prediction**” to guess which way branch will go earlier in the pipeline.
 - If branch has been seen before, predict as what you did the last time.
 - If branch has not been seen before, assume forward branches are not taken and backward branches are taken.
 - Update state on predictor with results of branch when it is finally calculated.
 - **Only flush pipeline if branch prediction was incorrect.**